

Correctness of Bellman-Ford:

let $\text{len}(i, v)$ denote the length of a shortest path from s to v that contains at most i edges. Since the sh. path from s to v contains $\leq n-1$ edges, it's sufficient to prove that after i iterations $\text{len}(v) \leq \text{len}(i, v)$

By induction on i

Base case: $i=0$ $\text{len}(s) = 0 \leq \text{len}(0, s) = 0$
 $\text{len}(v \neq s) = +\infty = \text{len}(0, v \neq s)$

Inductive hypothesis: $\text{len}(v) \leq \text{len}(k, v) \quad \forall 1 \leq k < i$

Take $i \geq 1$ and a shortest path from s to v with $\leq i$ edges. Let (u, v) be the last edge of this path. Then

$$\text{len}(i, v) = w(u, v) + \underbrace{\text{len}(i-1, u)}$$



By the ind. hyp. $\text{len}(v) \leq \text{len}(i-1, v)$
In the i -th iteration we update

$$\begin{aligned}\text{len}(v) &= \min \{ \text{len}(v), \underbrace{\text{len}(u) + w(u, v)}_{\leq \text{len}(i-1, u) + w(u, v)} \} \\ &= \text{len}(i, v) \\ &\leq \text{len}(i, v) \quad \text{as desired}\end{aligned}$$

All-Pairs Shortest Paths (APSP)

input: a directed, weighted graph $G = (V, E)$

output: one of the following

- a) $\text{dist}(u, v)$ for every ordered vertex pairs
- b) a declaration that G contains a negative cycle

Obvious solution: invoke B-F once for every vertex

$\rightarrow O(m \cdot n^2)$ very high


Can we do better?

Yes, using dynamic programming

Outline :

- 1) the B-F algorithm has a dynamic programming formulation
- 2) (we won't see) one can adapt that formulation to APSP, obtaining an $O(mn^2)$ algorithm; an improved formulation can be made to run in $O(n^3 \log n)$
- 3) a different dyn. prog. strategy gives an $O(n^3)$ algorithm

Bellman-Ford via dynamic programming
what are the subproblems here?



obs.: P' is a shortest path (to a \neq destination) with fewer edges than P

then P' can be interpreted as a solution to a smaller subproblem

\Rightarrow idea: introduce a parameter i that restricts the maximum number of edges allowed in a path, with smaller subproblems having smaller edge budgets

\hookrightarrow serves as a measure of subproblem size

Subproblems: compute $\text{len}(i, v)$, the length of a shortest path from s to v that contains at most i edges. (If no such path exists, define $\text{len}(i, v)$ as $+\infty$.)
 $\longrightarrow O(n^2)$ subproblems

Obs: every subproblem works with the full input; the idea is to control the allowable size of the output.

Bellman-Ford recurrence

$$\text{len}(i, v) = \begin{cases} 0 & i=0 \text{ and } v=s \\ +\infty & i=0 \text{ and } v \neq s \\ \min \left\{ \begin{array}{l} \text{len}(i-1, v) \\ \min_{(u,v) \in E} \{ \text{len}(i-1, u) + w(u, v) \} \end{array} \right\} & \text{otherwise} \end{cases}$$

(It's easy to transform a dyn.-prog. evaluation of this recurrence into our original formulation of Bellman-Ford)

This formulation can be adapted to APSP $\rightarrow O(n^3 \log n)$

The Floyd-Warshall algorithm

Idea: go one step further: instead of restricting the number of edges allowed in a solution, restrict the identities of the vertices that are allowed in a solution. (In other words, now path can pass through only certain vertices)

Let's define the subproblems:

- call the vertices $1, 2, \dots, n$


compute $\text{dist}(u, v, k) = \text{length of a shortest path from } u \text{ to } v \text{ that uses only vertices from } \{1, 2, \dots, k\} \text{ as internal (i.e., not } u \text{ or } v) \text{ vertices, and that does not contain a directed cycle. (If no such path exists, define } \text{dist}(u, v, k) \text{ as } +\infty.)$

$k \rightarrow$ measures the subproblem size

$\rightarrow O(n^3)$ subproblems

Algorithm: expand the set of allowed internal vertices, one vertex at a time, until this set is V .

Payoff of defining subproblems in this way: only 2 candidates for the optimal solution to a subproblem, depending on whether it uses vertex k or not:



$\Rightarrow O(1)$ work per subproblem

\Rightarrow Complexity : $O(n^3)$

Floyd-Warshall (G)

label the vertices $V = \{1, 2, \dots, n\}$ arbitrarily

// subproblems (k indexed from 0)

$A = n \times n \times (n+1)$ array

// base cases ($k=0$)

for $u=1$ to n do

for $v=1$ to n do

if $u=v$ then $A[u, v, 0] = 0$

else if $(u, v) \in E$ then $A[u, v, 0] = w(u, v)$

else $A[u, v, 0] = +\infty$

// solve all subproblems

for $k=1$ to n do

for $u=1$ to n do

for $v=1$ to n do

$A[u, v, k] = \min \{ A[u, v, k-1],$

$A[u, k, k-1] + A[k, v, k-1] \}$

\\ check for a negative cycle

for $v=1$ to n do

if $A[v, v, n] < 0$ then

return "G contains a negative cycle"

Is there a truly-subcubic algorithm for APSP?

↓

$$O(n^{3-\epsilon})$$

for some constant
 $\epsilon > 0$