


Example:

A



Correctness: follows from the Theorem

Complexity: $O(m \cdot n)$
(assume G represented
with an adjacency list)

polynomial time
 \Rightarrow efficient

Is $O(m \cdot n)$ really efficient in practice?

Think of Facebook graph: $n \simeq 2B$
 $m \simeq 2B \cdot \text{hundreds}$

So, not so efficient in very large graphs

Key observation: in the basic implementation the computation of the light edge is done repeatedly \Rightarrow should speed it up


golden rule in algorithms/coding: when an alg. repeats frequently the same operation, look for the "right" data structure to speed that operation up

\rightarrow here, a priority queue is what we need

\rightarrow implemented with a Heap:

Recap: {
INSERT: add an object to the heap
EXTRACT-MIN: remove an object with the smallest key
DELETE: given a pointer to an object, remove it

\rightarrow in a heap with n objects: $O(\log n)$



it's simpler to store vertices in the heap
(instead of edges)

Prim's implementation with a heap:

Prim(G, s)

for each $v \in V$ do
 $key(v) = +\infty$

// min. weight of any edge
 connecting v to the tree

$\pi(v) = \text{NULL}$

// parent of v in the tree

$key(s) = 0$

$H = V$ // contains all vertices not in tree
while $H \neq \emptyset$ do

$v^* = \text{extractMin}(H)$

for each v adjacent to v^* do

// update Key and π of each vertex
adjacent to v^* but not in tree

if $v \in H$ and $w(v^*, v) < \text{Key}(v)$

$\pi(v) = v^*$

// $\text{Key}(v) = w(v^*, v)$

delete v from H

$\text{Key}(v) = w(v^*, v)$

insert v into H

$$A = \{ (v, \pi(v)) : v \in V \setminus \{s\} \setminus \{H\} \}$$

Complexity: init $\rightarrow O(n)$

while $\rightarrow n$ iterations

extractMin $\rightarrow O(\log n)$

total cost of extractMin $O(n \log n)$

for loop : executed $O(m)$ times
in total

- $v \in H \rightarrow O(1)$

- $\text{key}(v) \rightarrow \text{delete+insert} : O(\log n)$

total cost of for loop $O(m \log n)$

Total : $O(n \log n + m \log n) =$

$= O(m \log n)$

↓
note: G is connected

it's near-linear time

Exercise : (Uniqueness of MSTs) Show that if the weights of the edges are all distinct then there exists exactly one MST.
(Hint : cut & paste argument)

Kruskal's algorithm (1956)

- it's very simple, very famous
- as fast as Prim, both in theory and in practice
- it gives us the opportunity to study a new data structure

GENERIC-NST(G) \rightarrow A : is a forest
 \rightarrow safe edge: a light edge connecting 2 distinct components

KRUSKAL(G) \parallel no more vertex needed

$$A = \emptyset$$

sort edges of G by weight


for each edge e , in nondecreasing order of weight do

if $A \cup \{e\}$ is acyclic then

$$A = A \cup \{e\}$$

return A


Example :



A

(simple optimization : stop the for loop when A has $n-1$ edges)

Correctness : follows from correctness of GENERIC-TST



Complexity: sorting: $O(m \log n)$

for loop: check whether $e=(u,v)$
closes a cycle is equivalent

to check whether A contains an
 $U-V$ path \rightarrow DFS on $G=(V,A)$
 \rightarrow complexity $O(n)$

Total: $O(m \cdot n)$

Can we implement it faster?