

Maximum Flows

Originated in the 50s to study rail networks

Definitions:

A flow network is a directed graph $G=(V,E)$ where each edge has a capacity $c(e) \in \mathbb{R}^+$, along with a designated source $s \in V$ and sink $t \in V$.

For convenience, write $c(e)=0$ if $e \notin E$, no edges enter s and no edges leave t .

A flow is a function $f: E \rightarrow \mathbb{R}^+$ satisfying the following constraints:


- 1) (capacity) $\forall e \in E \quad f(e) \leq c(e)$
- 2) (conservation) $\forall u \in V \setminus \{s, t\}$ we have

$$\sum_{\substack{v \in V \\ \text{s.t. } (v,u) \in E}} f(v,u) = \sum_{\substack{v \in V \\ \text{s.t. } (u,v) \in E}} f(u,v)$$

The value of a flow is $|f| = \sum_{v \in V} f(s, v)$
s.t. $(s, v) \in E$

Maximum flow problem: given a flow network
find a flow f of maximum value.

Example:



a flow of value 10


Applications: rail/airline/road networks, electrical networks, communication networks, liquid transportation networks; moreover, it can be applied to solve several other problems in computer science (e.g. bipartite matching)

Max flow reduces to linear programming (like many other problems) but there are more efficient special-purpose algorithms. We'll see one (Ford-Fulkerson) but there are plenty more efficient algorithms \rightarrow see Further Reading.

A natural idea: be greedy:

- find a path from s to t (in linear time with BFS)
- send as much flow along it as possible
- update capacities
- remove edges that have 0 remaining capacity
- repeat until the graph has no s - t paths

This won't always work:



Idea: revise/undo some of this flow later in the algorithm; how? By "pushing back" some flow through new edges in the reverse direction.

Definition: given a flow network G and a flow f , the residual network of G with respect to flow f , G_f , is a network with vertex set V and with edge set E_f as follows:

for every edge $e = (u, v)$ in G

— if $f(e) < c(e)$, add e to G_f with capacity

$$c_f(e) = c(e) - f(e)$$

— if $f(e) > 0$, add another edge (v, u) to G_f with capacity

$$c_f(e) = f(e)$$

The Ford-Fulkerson algorithm repeatedly finds an s - t path P in G_f (e.g. using BFS) and uses P to increase the current flow. P is called augmenting path.

Ford-Fulkerson (G, s, t) 1956

initialize $f(e) = 0$ for all $e \in G.E$

$G_f = G$

while there exists an augmenting path P in G_f do

let $\Delta_P = \min_{e \in P} c_f(e)$ Δ_P is the "bottleneck" capacity in P

for each edge $e = (u, v) \in P$ do

if $(u, v) \in G.E$ then

$f(u, v) = f(u, v) + \Delta_P$

else $f(v, u) = f(v, u) - \Delta_P$


update the residual graph G_f

return f




Example:




residual graph
↓




augmenting path



new flow
↓







no augmenting paths \Rightarrow

is a max flow,
of value 23


Complexity: assuming capacities are integers;
then

— the flow value increases by ≥ 1
in each iteration

— the complexity of each iteration
is $O(m)$

$$\Rightarrow O(m \cdot |f^*|) \quad f^* \text{ is a max flow}$$

A flow network for which F-F can take
 $\Theta(m \cdot |f^*|)$ time:



input size : $O(m \log U)$ $U = \text{max capacity}$

$O(m |f^*|) = O(m n U)$ "pseudo-polynomial"